

```
In [1]: from pykalman import KalmanFilter
import numpy as np
import pandas as pd
import sys
import matplotlib
import matplotlib.pyplot as plt
from skimage.color import lab2rgb
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
import skimage
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import FunctionTransformer
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from functools import reduce
import statsmodels.api as sm
lowess = sm.nonparametric.lowess
from scipy import stats
```

```
In [2]: def to_timestamp(dateTime):
        return dateTime.timestamp()

def map_genre(row):
    result = []
    for genre_code in row:
        matches = genres[genres['wikidata_id'] == genre_code]['genre_label'].values
        for match in matches:
            result.append(match)
    return result
```

```
In [3]: wikidata = pd.read_json('movies/data/wikidata-movies.json.gz', orient='record', lines=True)
#wikidata = pd.read_json('movies/data/wikidata-movies.json.gz', orient='record', lines=True)
genres = pd.read_json('movies/data/genres.json.gz', orient='record', lines=True)
```

```
In [4]: wikidata = wikidata[wikidata['made_profit'].notnull()].reset_index(drop=True)
```

```
In [5]: #movies = movies.copy() #from https://stackoverflow.com/questions/31468176/setting-index-to-column
#movies['genre_names'] = movies.apply(map_genre, axis=1)
wikidata['genre_names'] = wikidata['genre'].apply(map_genre)
wikidata['publication_timestamp'] = wikidata['publication_date'].apply(to_timestamp)
```

```
In [6]: rotten_tomatoes = pd.read_json('movies/data/rotten-tomatoes.json.gz', orient='record', lines=True)
```

```
In [7]: #rotten_tomatoes
rotten_tomatoes.columns
```

```
Out[7]: Index(['audience_average', 'audience_percent', 'audience_ratings',
              'critic_average', 'critic_percent', 'imdb_id', 'rotten_tomatoes_id'],
              dtype='object')
```

```
In [8]: omdb = pd.read_json('movies/data/omdb-data.json.gz', orient='record', lines=True)
```

```
In [9]: #omdb
```


```
In [10]: combined = wikidata.join(rotten_tomatoes.set_index('rotten_tomatoes_id'), on='rotten_tomatoes_id')
```

```
In [11]: combined = combined.join(omdb.set_index('imdb_id'), on='imdb_id')
```

```
In [12]: combined
```

1.0	...	1.278547e+09	4.2	91.0	568239.0	8.1	86
1.0	...	1.358381e+09	3.4	55.0	143566.0	6.0	65
1.0	...	6.946560e+07	4.4	98.0	731426.0	9.3	98
1.0	...	1.325376e+09	4.2	90.0	207900.0	8.4	96

```
In [13]: plt.title('Popularity Distribution')
plt.xlabel('Rank')
plt.ylabel('Views')
plt.scatter(combined['critic_average'], combined['audience_average'] * 2)
plt.show()
```



```
In [14]: test3 = combined[combined['audience_average'].notnull() & combined['critic_average'].notnull()]
print(stats.normaltest(test3['audience_average']).pvalue) #<0.05, therefore not normal
print(stats.mannwhitneyu(test3['critic_average'], test3['audience_average'] * 2).pvalue)

0.0012723373652919845
8.904724420410354e-15
```

```
In [15]: chi2, p, dof, expected = stats.chi2_contingency([test3['critic_average'].values, test3['audience_average'].values])
print(p) #>0.05, therefore one has no effect on the other?
print(expected)

1.0
[[5.5546799  7.87989474 7.16941243 ... 5.9422157  7.87989474 8.39660915]
 [3.0453201  4.32010526 3.93058757 ... 3.2577843  4.32010526 4.60339085]]
```

```
In [16]: # chi2, p, dof, expected = stats.chi2_contingency([test3['critic_average'].values, test3['audience_average'].values])
# print(p) #>0.05, therefore one has no effect on the other?
# print(expected)
```


```
In [17]: #combined.groupby('genre_names')
#pd.value_counts(combined.groupby('genre_names'), sort=False)
## TODO: Count distribution of genres and graph on histogram
```

Have average ratings changed over time?


```
In [18]: critic_average_test = combined[['publication_date', 'publication_timestamp', 'critic_average']]
fit = stats.linregress(critic_average_test['publication_timestamp'], critic_average_test['critic_average'])
critic_average_test['prediction'] = critic_average_test['publication_timestamp'] * fit[0] + fit[1]
print(fit.pvalue) #p < 0.05, therefore we can conclude that critic ratings are decreasing over time

6.156831173958255e-08
```

```
In [19]: plt.plot(critic_average_test['publication_date'], critic_average_test['critic_average'])
plt.plot(critic_average_test['publication_date'], critic_average_test['prediction'])
plt.show()
```




```
In [20]: plt.hist(np.subtract(critic_average_test['critic_average'],critic_average_test['p
plt.show()
#This is close enough to being normal.
#We expect a greater decline on the high end because the average critic rating is
```




```
In [21]: audience_average_test = combined[['publication_date','publication_timestamp','aud
fit = stats.linregress(audience_average_test['publication_timestamp'], audience_a
audience_average_test['prediction'] = audience_average_test['publication_timestam
print(fit.pvalue) #p > 0.05, therefore we cannot conclude that the audience rating
0.20019655801512012
```

```
In [22]: plt.plot(audience_average_test['publication_date'], audience_average_test['audien
plt.plot(audience_average_test['publication_date'], audience_average_test['predic
plt.show()
```



```
In [23]: plt.hist(np.subtract(audience_average_test['audience_average'], audience_average_t
plt.show())
#This is close enough to being normal.
#We expect a greater decline on the high end because the average audience rating
```




Do average audience ratings change based on its popularity?


```
In [24]: audience_ratings_test = combined[['publication_date', 'publication_timestamp', 'aud
#Removing movies with n >= 10000000 ratings as they seem like outliers
audience_ratings_test = audience_ratings_test[audience_ratings_test['audience_rat
fit = stats.linregress(audience_ratings_test['audience_ratings'], audience_rating
audience_ratings_test['prediction'] = audience_ratings_test['audience_ratings']*f
print(fit.pvalue) #p < 0.05, therefore we can conclude that higher averages corre
```

0.000308565324974134

```
In [25]: plt.plot(audience_ratings_test['audience_ratings'], audience_ratings_test['audien
plt.plot(audience_ratings_test['audience_ratings'], audience_ratings_test['predic
plt.show()
```



```
In [26]: plt.hist(np.subtract(audience_ratings_test['audience_average'], audience_ratings_t
plt.show()
#This is close enough to being normal.
#We expect a greater decline on the high end because the average audience rating
```



Does genre have an effect on profitability?

```
In [27]: def genre_agg(combined_row):
    for genre_id in combined_row['genre']:
        genre_test.loc[genre_test['wikidata_id'] == genre_id, 'total'] += 1
        if (combined_row['made_profit'] == 1.0):
            genre_test.loc[genre_test['wikidata_id'] == genre_id, 'profit'] += 1
```

```
In [28]: genre_test = genres
genre_test['profit'] = 0
genre_test['total'] = 0
combined.apply(genre_agg, axis=1)
genre_test = genre_test[genre_test['total'] > 0]
```

```
In [29]: genre_test['loss'] = genre_test['total'] - genre_test['profit']
contingency = genre_test[['profit', 'loss']]
contingency = contingency[contingency['profit'] >= 5]
contingency = contingency[contingency['loss'] >= 5]
chi2, p, dof, expected = stats.chi2_contingency(contingency)
print(p) # p < 0.05, therefore genre has some effect on profitability
```

0.01956332775267009

C:\Users\User\Anaconda3\lib\site-packages\ipykernel_launcher.py:1: SettingWithC
opyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy> (<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)

"""Entry point for launching an IPython kernel.

In []:

```
In [30]: #NLP
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfTransformer
```

```
In [31]: count_vect = CountVectorizer()
X_train_counts = count_vect.fit_transform(test3["omdb_plot"])
```

```
In [51]: tfidf_transformer = TfidfTransformer()
X_train_tfidf = tfidf_transformer.fit_transform(X_train_counts)
```

```
In [52]: from sklearn.naive_bayes import MultinomialNB
y=test3["audience_average"]
y=y.astype('int')
clf = MultinomialNB().fit(X_train_tfidf, y)
```

```
In [62]: docs_new = ['love', 'Fat']
X_new_counts = count_vect.transform(docs_new)
X_new_tfidf = tfidf_transformer.transform(X_new_counts)
predicted = clf.predict(X_new_tfidf)
```

```
In [63]: for doc, category in zip(docs_new, predicted):
    print('%r => %s' % (doc, test3["omdb_plot"][category])) #not working?
```

'love' => Lt. John Dunbar is dubbed a hero after he accidentally leads Union troops to a victory during the Civil War. He requests a position on the western frontier, but finds it deserted. He soon finds out he is not alone, but meets a wolf he dubs "Two-socks" and a curious Indian tribe. Dunbar quickly makes friends with the tribe, and discovers a white woman who was raised by the Indians. He gradually earns the respect of these native people, and sheds his white-man's ways.

'Fat' => Lt. John Dunbar is dubbed a hero after he accidentally leads Union troops to a victory during the Civil War. He requests a position on the western frontier, but finds it deserted. He soon finds out he is not alone, but meets a wolf he dubs "Two-socks" and a curious Indian tribe. Dunbar quickly makes friends with the tribe, and discovers a white woman who was raised by the Indians. He gradually earns the respect of these native people, and sheds his white-man's ways.